TRACK 1 · MULTIMODAL CONTENT CREATION

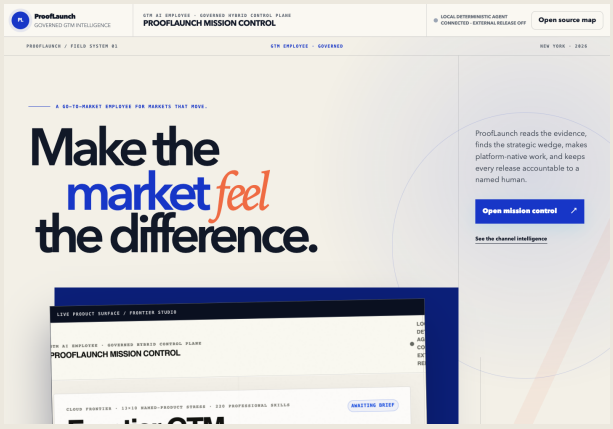
ProofLaunch

Multimodal GTM Studio for AMD Radeon / ROCm

TEAM BANYAN

INPUT	Owned image / PDF / audio / video / text evidence
OUTPUT	Audited image + video campaign pack + evidence ledger
CONTROL	One Supervisor · human approval · maximum two repairs
STATUS	Product, video and PR ready; fresh end-to-end Radeon run pending

From product evidence to publish-safe campaign assets—without fabricated claims, redrawn products, or hidden repair history.



01 · PROJECT BACKGROUND

Creative speed is not enough when product truth breaks.

Generative campaign tools can produce many variants quickly, but unsupported claims, altered product identity, unreadable text, unsafe platform layouts, and missing provenance turn generation speed into release risk.

ProofLaunch changes the unit of work

• Every commercial claim starts as an approved source span and Claim ID.

• The original product remains a protected compositing layer.

• Every candidate is audited before it can reach export.

• Repairs are explicit, bounded, and retained in the evidence history.

• Unknown or unavailable proof blocks release instead of becoming a claim.

WITHOUT PROOFLAUNCH	WITH PROOFLAUNCH
Prompt → image → manual guess	Evidence → gated campaign pack
Untraceable text claims	Source span → Claim ID → exact copy
Product may be redrawn	Identity Lock protects the original layer
Silent retries	Maximum two recorded repair rounds
Export by appearance	Export by inspectable release gates

Built for the operator who must defend the final asset.

Primary users: campaign operators, owner-reviewers, and technical qualifiers at small ecommerce and hardware organizations.

Primary scenario

01 EVIDENCE
Hash + rights
Located claims

02 PLAN
Claim-bound
channel intent

03 CREATE
Radeon-targeted
media candidates

04 AUDIT
Six release
gates

05 REPAIR
≤2 bounded
rounds

06 DELIVER
Human gate
+ hash ledger

HUMAN DECISION	AUTOMATED ASSISTANCE
Approve source claims	Extract candidate source spans
Approve risky repairs	Produce bounded repair proposal
Approve external side effects	Inspect and prepare deliverables
Accept final release evidence	Hash and package evidence

One Supervisor; many bounded capabilities.

The Supervisor is the only component allowed to change campaign state, consume retry budgets, enter export, or request approval for a side effect. External agents return observations or typed intents.

PROOFLAUNCH SUPERVISOR	Owns state, approvals, retry budgets and export
EVIDENCE + PLAN	Qwen3-VL · PDF/Whisper · Claim Ledger · Director
CREATE	Radeon-targeted SDXL / LTX adapters · deterministic typography
AUDIT + REPAIR	SigLIP2 · GroundingDINO / SAM2 · CLAP · bounded repair
HUMAN-GATED DELIVERY	Campaign pack · SHA-256 evidence ledger · no autonomous publish/send/spend

Agent remix boundary

• Diffusers supplies the actual optional local Radeon generation adapter.

• Pillow supplies deterministic typography and immutable-layer compositing.

• ComfyUI remains an AMD-supported alternative backend, not bundled code.

• Qwen3-VL, GroundingDINO, Gemma, SDXL, SigLIP2, LTX, SAM2, and CLAP are configured for sequential residency.

• No upstream framework owns ProofLaunch release state.

Models propose; deterministic gates decide release.

CAPABILITY	IMPLEMENTATION	BOUNDARY
Input understanding	Qwen3-VL + GroundingDINO + PDF + Whisper	Located source evidence
Audience Pulse	Gemma 4 E4B-it + six-label emotion LoRA	Not conversion prediction or diagnosis
Radeon Studio	Radeon-targeted SDXL still + LTX video adapters	No closed online API on the key inference path
Identity Lock	Exact layer + decoded H.264 tolerance	Pre/post-codec identity
Video consistency	SAM2 tracking + CLAP + A/V timing	Temporal product and cue evidence
Typography	Deterministic rendering	Exact claims, model IDs, price, CTA
Release Audit	Mechanical + SigLIP2 score	Six fail-closed release gates
Evidence	SHA-256-linked JSONL ledger	Tampering invalidates verification

Fail-closed decisions

• No approved source-bound claim → stop.

• No generated asset → stop.

• Critical finding after two repairs → stop.

• No deliverable or provenance evidence → stop.

The GPU path is designed around evidence and bounded memory.

- Multi-GPU-first execution contract with a verified single-card fallback profile.
- No default bitsandbytes 4-bit dependency.
- ModelScope-local model and dataset snapshots.
- PyTorch CUDA-compatible API surface on ROCm.
- Sequential model loading and explicit release between large stages.
- Frozen 450-case matrix + 30 full audio-video cases.
- Protocol, V7 build, and named baseline settings are hash-bound.
- Nine phase-batched passes enforce one resident model at a time.
- Latency, peak VRAM, model reference, seed, and output hash are mandatory.

CURRENT V10 EVIDENCE	STATUS
Web product surface · 46 routes / 60 API contracts	Observed locally
UI / route validation / final integrity locks	17/17 · 10/10 · 4/4
Source/model/data drift locks	Observed locally
Radeon execution contract	Mechanically verified
Historical Gemma-4 Radeon / HIP execution	Observed
Fresh end-to-end ProofLaunch Radeon run	Not observed
Real post-training model evaluation	Not observed
Real image/video generation benchmark	Not observed
Independent 4A human evaluation	Not observed
External publication or spend	Disabled

Submission readiness is machine-readable and fail-closed.

EVIDENCE	RESULT
Product-critical regression	827 / 827 passed
UI tests	17 / 17 passed
Route validation	10 / 10 passed
Web product surface	46 routes / 60 API contracts
Final integrity locks	4 / 4 passed
Run-bound deliverables	10 / 10 registered
Automatic publish / send / spend	0
Source/model/data lock drift	Fail-closed
Historical Gemma-4 Radeon receipt	Observed
Fresh end-to-end Radeon campaign receipt	Not observed
Demo Video · 190.9 seconds	Included
Official fork + compliant PR	Open · #308
Final Evidence Seal	Not sealed

Official Track 1 package

• English Project Profile PDF.

• Complete source repository and reproduction README.

• Demo Video (190.9 seconds; official guidance recommends 3–5 minutes).

• Supplementary English deck.

• Fork + English PR titled: Track 1, Banyan, ProofLaunch: Multimodal GTM Studio.

Historical Gemma-4 inference and LoRA training on Radeon/HIP are preserved. This does not replace the fresh end-to-end ProofLaunch campaign run, latency/VRAM capture, or final output hashes required by the remaining P0 gate.