SUMBOT

Local project materials in. Source-grounded weekly report out.

THE PROBLEM

Weekly reporting is repetitive, while generic summarization can overstate progress, lose evidence, or confuse old plans with current results.

THE SOLUTION

Sumbot uses a fixed Extractor -> Writer -> Reviewer workflow so every stage stays inspectable and every report remains subject to human approval.

WHY LOCAL

An OpenAI-compatible endpoint keeps the model-serving layer replaceable and enables controlled deployment with AMD Radeon acceleration in the inference backend.

1 LOAD	2 EXTRACT	3 WRITE	4 REVIEW	5 EXPORT
Paste text or upload supported project files	Collect current facts with source filenames	Draft a compact report from extracted facts only	Check claims against facts and original materials	Human review, copy, and Markdown download

BUILT FOR AN INSPECTABLE LOCAL WORKFLOW

- Multi-format input: text, Word, PowerPoint, Excel, and PDF
- File-level source attribution in facts and exports
- Background queue and local SQLite persistence
- Stage timings, warnings, and explicit failure states
- Deterministic mock mode for UI validation

AMD RADEON DEPLOYMENT

Sumbot connects to a local OpenAI-compatible model server. The backend owns ROCm kernels, model format, quantization, and GPU memory settings; Sumbot reduces contention with one model job at a time and bounded stage outputs.

VALIDATION BOUNDARY

No unverified performance number is claimed. The final demo must show the actual GPU, runtime, model, MOCK_LLM=false, and observed timings. Mock output is not GPU evidence.